

Puppet Tutorial

By Alessandro Franceschi

Example42

11-07-2015

Puppet Essentials -

Overview

Introduction to Puppet

Configuration management tools

Puppet Ecosystem and related software

What is Puppet

A Configuration Management Tool

A framework for Systems Automation

A Declarative Domain Specific Language (**DSL**)

An OpenSource software in written Ruby

Works on Linux, Unix (Solaris, AIX, *BSD), MacOS, Windows (Supported Platforms)

Developed by Puppet Labs

Used by <http://puppetlabs.com/customers/companies/>

...

... and many others

Configuration Management advantages

Infrastructure as Code: Track, Test, Deploy, Reproduce, Scale

Code commits log shows the **history of change** on the infrastructure

Reproducible setups: Do once, repeat forever

Scale quickly: Done for one, use on many

Coherent and consistent server setups

Aligned Environments for devel, test, qa, prod nodes

Alternatives to Puppet: Chef, CFEngine, Salt, Ansible

References and

Ecosystem

Puppet Labs - The Company behind Puppet

Puppet - The OpenSource version

Puppet Enterprise - The commercial version

The Community - Active and vibrant

Puppet Documentation - Main and Official reference

Puppet Modules on: Module Forge and GitHub

Software related to Puppet:

Facter - Complementary tool to retrieve system's data

MCollective - Infrastructure Orchestration framework

Hiera - Key-value lookup tool where Puppet data can be placed

PuppetDB - Stores all the data generated by Puppet

Puppet Dashboard - A Puppet *Web Frontend* and External Node Classifier (ENC)

The Foreman - A well-known third party provisioning tool and Puppet ENC

Gepetto - A Puppet IDE based on Eclipse

Installation

Debian, Ubuntu
Available by default

```
apt-
get install puppet      # On clients (nodes)
apt-
get install puppetmaster # On server (master)
```

RedHat, Centos, Fedora
Add EPEL repository or RHN Extra channel

```
yum install puppet      # On clients (nodes)
yum install puppet-
server # On server (master)
```

Use PuppetLabs repositories for latest updates

Installation Instructions for different OS

Puppet Versions

From version 2 Puppet follows Semantic Versioning standards to manage versioning, with a pattern like:
MAJOR.MINOR.PATCH

This implies that MAJOR versions might not be backwards compatible, MINOR versions may introduce new features keeping compatibility and PATCHes are for backwards compatible bug fixes.

This is the history of the main Puppet versions and some relevant changes
Check also [Puppet Language History](#) for a more complete review:

0.9.x - First public beta

0.10.x - 0.21.x - Various "early" versions

0.22.x - 0.25.x - Improvements and OS wider support

2.6.x - Many changes. Parametrized classes

2.7.x - Windows support

3.0.x - Many changes. Disabled dynamic variables scope. Data bindings

3.2.x - Future parser (experimental, will be default in Puppet 4)

4.x.x - Release in early 2015. New parser, new type system and lot of changes, some backwards incompatibilities

Language Basics - Overview

**Introduction to Puppet language
and terms**

Resource types

Classes and defines

Variables and parameters

Nodes classification

Puppet Language

A **Declarative** Domain Specific Language (DSL)

It defines **STATES** (Not procedures)

Puppet code is written in **manifests** (files with **.pp** extension)

In the code we declare **resources** that affect elements of the system (files, packages, services ...)

Resources are grouped in **classes** which may expose parameters that affect their behavior.

Classes and configuration files are organized in **modules**.

Consult the **official glossary** to give the correct meaning to Puppet terms

Nodes classification

When clients connect, the Puppet Master generates a

catalog with the list of resources that clients have to apply locally.

The Puppet Master has to *classify* nodes and define for each of them:

The **classes** to include

The **parameters** to pass

The Puppet environment to use

The **catalog** is generated by the Master according to the logic of our Puppet code and data.

In our code we can define our **variables** and use other ones that may come from different sources:

facts generated directly by the client

parameters obtained from node's classification

Puppet **internal** variables

Resource Types (Types)

Resource Types are single **units of configuration** composed by:

A **type** (package, service, file, user, mount, exec ...)

A **title** (how is called and referred)

Zero or more **arguments**

```
type { 'title':  
    argument => value,  
    other_arg => value,  
}
```

Example for a **file** resource type:

```
file { 'motd':  
    path    => '/etc/motd',  
    content => 'Tomorrow is another day',  
}
```

Resource Types reference

Find online the complete Type Reference for the latest or earlier versions.

From the shell the command line interface:

```
puppet describe file
```

For the full list of available descriptions try:

```
puppet describe --list
```

Give a glance to [Puppet code](#) for the list of **native** resource types:

```
ls $(facter rubysitedir)/puppet/type
```

Simple samples of resources

Installation of OpenSSH package

```
package { 'openssh':
  ensure => present,
}
```

Creation of /etc/motd file

```
file { 'motd':
  path => '/etc/motd',
}
```

Start of httpd service

```
service { 'httpd':
  ensure => running,
  enable => true,
}
```

More complex samples of resources

Management of nginx service with parameters defined in module's variables

```
service { 'nginx':
  ensure      => $::nginx::manage_service
  _ensure,
  name        => $::nginx::service_name,
  enable      => $::nginx::manage_service
  _enable,
}
```

Creation of nginx.conf with content retrieved from different sources (first found is served)

```
file { 'nginx.conf':
  ensure  => present,
  path    => '/etc/nginx/nginx.conf',
  source  => [
    "puppet:///modules/site/nginx.conf
--${::fqdn}",
    "puppet:///modules/site/nginx.conf
" ],
}
```

Resource Abstraction Layer

Resources are abstracted from the underlying OS

Resource **types** have different providers for different OS

The **package** type is known for the great number of providers

```
ls $(facter rubysitedir)/puppet/provider  
/package
```

Use **puppet resource** to interrogate the RAL:

```
puppet resource user  
puppet resource user root  
puppet resource package  
puppet resource service
```

Or to directly modify resources:

```
puppet resource service httpd ensure=run  
ning enable=true
```

Classes - Definition

Classes are **containers** of different resources. Since Puppet 2.6 they can have parameters

Example of a class **definition**:

```
class mysql (
    root_password => 'default_value',
    port          => '3306',
) {
    package { 'mysql-server':
        ensure => present,
    }
    service { 'mysql':
        ensure     => running,
    }
    [...]
}
```

Note that when we define a class we just describe what it does and what parameters it has, we don't actually add it and its resources to the catalog.

Classes - Declaration

When we have to use a class previously defined, we **declare** it.

This can be done in 2 different ways:

"Old style" class declaration, without parameters:

```
include mysql
```

(Inside a catalog we can have multiple includes of the same class but that class it's applied only once.)

"New style" (from [Puppet 2.6](#)) class declaration with explicit parameters:

```
class { 'mysql':
  root_password => 'my_value',
  port          => '3307',
}
```

(Syntax is the same of normal resources and the same class can be declared, in this way, only once inside the same catalog)

Defines

Also called: **Defined resource types** or **defined types**

Similar to parametrized classes but can be used multiple times (with different titles).

Definition of a define:

```
define apache::virtualhost (
  $ensure  = present,
  $template = 'apache/virtualhost.conf.e
rb' ,
  [...] ) {

  file { "ApacheVirtualHost_${name}":
    ensure  => $ensure,
    content => template("${template}"),
  }
}
```

Declaration of a define:

```
apache::virtualhost { 'www.example42.com
':
  template => 'site/apache/www.example42
.com-erb'
}
```

Variables

We need them to provide different configurations for different kind of servers.

They can be defined in different places and by different actors:

Can be provided by client nodes as **facts**

Can be **defined by users** in Puppet code, on Hiera or in the ENC

Can be **built-in** and be provided directly by Puppet

Facts

Facter runs on clients and collects **facts** that the server can use as variables

```
a1$ facter  
  
architecture => x86_64  
fqdn => Macante.example42.com  
hostname => Macante  
interfaces => lo0,eth0  
ipaddress => 10.42.42.98
```

```
ipaddress_eth0 => 10.42.42.98
kernel => Linux
macaddress => 20:c9:d0:44:61:57
macaddress_eth0 => 20:c9:d0:44:61:57
memorytotal => 16.00 GB
netmask => 255.255.255.0
operatingsystem => Centos
operatingsystemrelease => 6.3
osfamily => RedHat
virtual => physical
```

User Variables

We can define custom variables in different ways:

In Puppet manifests:

```
$role = 'mail'

$package = $::operatingsystem ? {
    /(?:
i:Ubuntu|Debian|Mint)/ => 'apache2',
    default                  => 'http
d',
}
```

In an External Node Classifier (ENC)

Commonly used ENC are Puppet DashBoard, the Foreman, Puppet Enterprise.

In an Hiera backend

```
$syslog_server = hiera(syslog_server)
```

Built-in variables

Puppet provides some useful built in variables, they can be:

Set by the client (agent)

`$clientcert` - the name of the node's certificate. By default its `$::fqdn`

`$clientversion` - the Puppet version on the client

Set by the server (master)

`$environment` (default: `production`) - the Puppet environment where are placed modules and manifests.
`$servername`, `$serverip` - the Puppet Master FQDN and IP

`$serverversion` - the Puppet version on the server
`$settings::<name>` - any configuration setting on the Master's `puppet.conf`

Set by the server during catalog compilation

`$module_name` - the name of the module that contains the current resource's definition
`$caller_module_name` - the name of the module that contains the current resource's declaration

Environments

Puppet environments allow isolation of Puppet code and data: for each environment we can have different paths manifest files, Hiera data and modules.

Puppet's environments DO NOT necessarily have to match the operational environments of our servers.

Environment management has changed from Puppet 3.6 onwards:

Earlier versions were based on so-called **Config file Environments** where each environment had to be

environments where each environment had to be defined in **puppet.conf**.

From 3.6 **directory environments** have been introduced and the older approach has been deprecated.

On Puppet 4.x only directory environments are supported.

Config file Environments

The "old" config file environments are defined inside **puppet.conf** with a syntax like:

```
[test]
  modulepath = $confdir/environments/test/modules:$confdir/modules
  manifest = $confdir/environments/test/manifests
```

For the default **production** environment there were the config parameters, now deprecated, **manifest** and **modulepath** in the **environment** section

mountpath in the main section.

Directory Environments

Directory Environments are configured in `puppet.conf` as follows:

```
[main]
environmentpath = $configdir/
```

Then inside the
`/etc/puppet/environments/$environment/`
directory we have:

```
modules/ # Directory containing module
s
manifests/ # Directory containing site.p
p
environment.conf # Conf file for the env
ironment
```

Nodes - Default classification

A node is identified by the PuppetMaster by its **certname**, which defaults to the node's **fqdn**

In the first manifest file parsed by the Master, `site.pp`, we can define nodes with a syntax like:

```
node 'web01' {
  include apache
}
```

We can also define a list of matching names:

```
node 'web01' , 'web02' , 'web03' {
  include apache
}
```

or use a regular expression:

```
node /^www\d+$/ {
  include apache
}
```

A node can inherit another node and include all the classes and variables defined for it, this feature is now deprecated and is not supported anymore on [Puppet 4](#).

Nodes classification via an ENC

Puppet can query an external source to retrieve the classes and the parameters to assign to a node. This source is called External Node Classifier (ENC) and can be anything that, when interrogated via a script with the clients' certname as first parameter, returns a yaml file with the list of classes and parameters.

Common ENC are **Puppet DashBoard**, **The Foreman** and **Puppet Enterprise** (where the functionality of ENC is enabled by default).

To enable the usage of an ENC set these parameters in `puppet.conf`

```
# Enable the usage of a script to classify nodes
node_terminus = exec

# Path of the script to execute to classify nodes
external_nodes = /etc/puppet/node.rb
```

Nodes classification with Hiera

Hiera provides a **hiera_include** function that allows the inclusion of classes as defined on Hiera. This is an approach that can be useful when there's massive usage of Hiera as backend for Puppet data.

In `/etc/puppet/manifests/site.pp` just place:

```
hiera_include('classes')
```

and place, as an array, the classes to include in our Hiera data source under the key `classes`.

The **catalog** is the complete list of resources, and their relationships, that the Puppet Master generates for the client.

It's the result of all the puppet code and logic that we define for a given node in our manifests and is applied on the client after it has been received from the master.

The client uses the RAL (Resource Abstraction Layer) to execute the actual system's commands that convert abstract resources like

```
package { 'openssh': }
```

to their actual fulfillment on the system (apt-get install openssh , yum install openssh ...).

The catalog is saved by the client in
`/var/lib/puppet/client_data/catalog/$certname.json`

Language Basics - Practice

On a test machine **install Puppet** if not already installed.

Practice with the **commands**:

puppet describe
puppet resource
facter

Create a **manifest** file and in that file manage the installation of the package and the management of the service of *nginx*.

Use `puppet apply` to apply the resources declared in our manifest.

Puppet Configuration and basic usage - Overview

Operational modes: apply / agent
Configuration files and options
Useful Paths

Anatomy of a Puppet run

Operational modes

Masterless - puppet apply

Our Puppet code (written in manifests) is applied directly on the target system.

No need of a complete client-server infrastructure.

Have to distribute manifests and modules to the managed nodes.

Command used: `puppet apply` (generally as root)

Master / Client - puppet agent

We have clients, our managed nodes, where Puppet client is installed.

And we have one or more Masters where Puppet server runs as a service

Client/Server communication is via https (**port 8140**)

Clients certificates have to be accepted (**signed**) on the Master

Command used on the client: `puppet agent` (generally as **root**)

Command used on the server: `puppet master` (generally as **puppet**)

Masterless setup

Puppet manifests are deployed directly on nodes and applied locally:

```
puppet apply --  
modulepath /modules/ /manifests/file.pp
```

More fine grained control on what goes in production for what nodes

Ability to trigger multiple truly parallel Puppet runs

No single point of failure, no Master performance issues

Need to define a fitting deployment workflow. Hints:
Rump - supply_drop

With Puppet > 2.6 we can use file sources urls like:

```
puppet:///modules/example42/apache/vhost  
.conf  
  
# they point to  
$modulepath/example42/apache/vhost.conf
```

StoreConfigs usage require access to Puppet DB from every node

Master / Client Setup

A Puppet server (running as 'puppet') listening on 8140 on the Puppet Master (the server)

A Puppet client (running as 'root') on each managed node

Client can be run as a service (default), via cron (with random delays), manually or via MCollective

Client and Server have to share SSL certificates. New client certificates must be signed by the Master CA

It's possible to enable automatic clients certificates signing on the Master (be aware of security concerns)

Certificates management

On the Master we can use `puppet cert` to manage certificates

List the (client) certificates to sign:

```
puppet cert list
```

List all certificates: signed (+), revoked (-), to sign ():

```
puppet cert list --all
```

Sign a client certificate:

```
puppet cert sign <certname>
```

Remove a client certificate:

```
puppet cert clean <certname>
```

Client stores its certificates and the server's public one in `$vardir/ssl**` (`/var/lib/puppet/ssl` on [Puppet OpenSource](#))

Server stores clients public certificates and in `$vardir/ssl/ca` (`/var/lib/puppet/ssl/ca`). DO NOT remove this directory.

Certificates management - First run

By default the first [Puppet](#) run on a client fails:

```
client # puppet agent -t
          # "Exiting; no certificate found
          and waitforcert is disabled"
```

An optional `--waitforcert 60` parameter makes
client wait 60 seconds before giving up.

The server has received the client's CSR which has to

The server has received the clients CSR which has to be manually signed:

```
server # puppet cert sign <certname>
```

Once signed on the Master, the client can connect and receive its catalog:

```
client # puppet agent -t
```

If we have issues with certificates (reinstalled client or other certs related problems):

Be sure client and server times are synced

Clean up the client certificate. On the client remove it:

```
client # mv /var/lib/puppet/ssl /var/lib/puppet/ssl.old
```

On the Master clean the old client certificate:

```
server # puppet cert clean <certname>
```

Puppet configuration: `puppet.conf`

It's Puppet main configuration file.
On opensource Puppet is generally in:

```
/etc/puppet/puppet.conf
```

On Puppet Enterprise:

```
/etc/puppetlabs/puppet/puppet.conf
```

When running as a normal user can be placed in the
home directory:

```
/home/user/.puppet/puppet.conf
```

Configurations are divided in [stanzas] for different
Puppet sub commands

Common for all commands: **[main]**

For puppet agent (client): **[agent]** (Was [puppetd] in
Puppet pre 2.6)

For puppet apply (client): **[user]** (Was [puppet])

For puppet master (server): **[master]** (Was
[puppetmasterd] and [puppetca])

Hash sign (#) can be used for comments.

Main configuration options

To view all or a specific configuration setting:

```
puppet config print all  
puppet config print modulepath
```

Important options under [main] section:

vardir: Path where Puppet stores dynamic data.

ssldir: Path where SSL certifications are stored.

Under [agent] section:

server: Host name of the PuppetMaster. (Default:

puppet)

certname: Certificate name used by the client.
(Default is its fqdn)

runinterval: Number of minutes between Puppet runs,
when running as service. (Default: 30)

report: If to send Puppet runs' reports to the
**report_server. (Default: true)

Under [master] section:

autosign: If new clients certificates are automatically
signed. (Default: false)

reports: How to manage clients' reports (Default:
store)

storeconfigs: If to enable store configs to support
exported resources. (Default: false)

Full configuration reference on the official site.

Common command-line parameters

All configuration options can be overridden by command-line options.

A very common option used when we want to see immediately the effect of a [Puppet run](#) (it's actually the combination of: `-onetime`, `-verbose`, `-ignorecache`, `-nodeamonize`, `-no-usecacheonfailure`, `-detailed-exit-codes`, `-no-splay`, and `-show_diff`):

```
puppet agent --  
test # Can be abbreviate to -t
```

Run puppet agent in foreground and debug mode:

```
puppet agent --test --debug
```

Run a dry-run puppet without making any change to the system:

```
puppet agent --test --noop
```

Run puppet using an environment different from the default one:

```
puppet agent --test --  
environment testing
```

Wait for certificate approval (by default 120 seconds) in the first [Puppet run](#) (useful during automated first time installation if `PuppetMaster.autosign` is `false`):

```
puppet agent --test --waitforcert 120
```

Useful paths

/var/log/puppet contains logs (but also on normal syslog files, with facility daemon), both for agents and master

/var/lib/puppet contains Puppet operational data (catalog, certs, backup of files...)

/var/lib/puppet/ssl contains SSL certificate

/var/lib/puppet/clientbucket contains backup copies of the files changed by Puppet

/etc/puppet/manifests/site.pp (On Master) The first manifest that the master parses when a client connects in order to produce the configuration to apply to it (Default on Puppet < 3.6 where are used **config-file environments**)

/etc/puppet/environments/production/manifests/site.pp (On Master) The first manifest that the master parses when using **directory environments** (recommended from Puppet 3.6 and default on Puppet >= 4)

/etc/puppet/modules and **/usr/share/puppet/modules** (On Master) The default directories where modules are searched

/etc/puppet/environments/production/modules (On Master) An extra place where modules are looked for when using **directory environments**

Other configuration

files:

auth.conf

Defines ACLs to access [Puppet's REST interface](#).

[Details](#)

fileserver.conf

Used to manage ACL on files served from sources different than modules [Details](#)

puppetdb.conf

Settings for connection to PuppetDB, if used. [Details](#)

tagmail.conf , autosign.conf , device.conf , routes.yaml

These are other configuration files for specific functions. [Details](#)

/etc/puppet/environments/production/environment

Contains environment specific settings

compilation

Execute Puppet on the client

Client shell # puppet agent -t

If `pluginsync = true` (default from Puppet 3.0) the client retrieves all extra plugins (facts, types and providers) present in modules on the server's `$modulepath`

Client output # Info: Retrieving plugin

The client runs facter and send its facts to the server

Client output # Info: Loading facts in
`/var/lib/puppet/lib/facter/...` [...]

The server looks for the client's hostname (or certname, if different from the hostname) and looks into its nodes list

The server compiles the catalog for the client using also client's facts

Server's logs # Compiled catalog for in environment production in 8.22 seconds

If there are not syntax errors in the processed Puppet code, the server sends the catalog to the client, in PSON format.

Anatomy of a Puppet Run - Part 2: Catalog application

Client output # Info: Caching catalog for

The client receives the catalog and starts to apply it locally

If there are dependency loops the catalog can't be applied and the whole run fails.

Client output # Info: Applying configuration version '1355353107'

All changes to the system are shown here. If there are errors (in red or pink, according to Puppet versions) they are relevant to specific resources but do not block the application of the other resources (unless they depend on the failed ones).

At the end of the Puppet run the client sends to the server a report of what has been changed

Client output # Finished catalog run in 13.78 seconds

The server eventually sends the report to a Report Collector

Puppet Configuration and basic usage - Practice

Use the command `puppet config print` to explore Puppet's configuration options.

Give a look to the various Puppet related directories and their contents:

`/etc/puppet`, `/var/lib/puppet`, `/var/log/puppet`

If the lab setup allows it, run Puppet on a new client and sign its certificate on the master.

Explore and describe the output of a Puppet run in apply and agent mode.

Resources reference - Overview

Where to find docs about Puppet resources

Main resources: package, service, file, user, exec

Language Style

Online documentation on types

Complete Type Reference

This is the complete reference for Puppet types based on the latest version.

Check here for the reference on all the older versions.

Check the Puppet Core Types Cheatsheet for an handy PDF with the essential reference.

Inline documentation on types

Types reference documentation

We can check all the types documentation via the command line:

```
puppet describe <type>
```

For a more compact output:

```
puppet describe -s <type>
```

To show all the available resource types:

```
puppet describe -l
```

Inspect existing resource types

To interactively inspect and modify our system's resources

```
puppet resource <type> [name]
```

Remember we can use the same command to CHANGE our resources attributes:

```
puppet resource <type> <name> [attribute =value] [attribute2=value2]
```

Managing packages

Installation of packages is managed by the **package** type.

The main arguments:

```
package { 'apache':
  name      => 'httpd',  # (namevar)
  ensure    => 'present' # Values: 'absent', 'latest', '2.2.1'
  provider  => undef,    # Force an explicit provider
}
```

Managing services

Management of services is via the **service** type.

The main arguments:

```
service { 'apache':
  name      => 'httpd',    # (namevar)
  ensure     => 'running'  # Values: 'stop
                           ped', 'running'
  enable     => true,       # Define if to
  enable service at boot (true|false)
  hasstatus  => true,       # Whether to us
  e the init script' status to check
                           # if the servic
  e is running.
  pattern   => 'httpd',    # Name of the p
  rocess to look for when hasstatus=false
}
```

Managing files

Files are the most configured resources on a system,
we manage them with the **file** type:

```
file { 'httpd.conf':
  # (namevar) The file path
  path      => '/etc/httpd/conf/httpd.
```

```
conf',
    # Define the file type and if it should exist:
    # 'present','absent','directory','link'
    ensure      => 'present',
    # Url from where to retrieve the file content
    source      => 'puppet://[puppetfileserver]/<share>/path',
    # Actual content of the file, alternative to source
    # Typically it contains a reference to the template function
    content     => 'My content',
    # Typical file's attributes
    owner       => 'root',
    group       => 'root',
    mode        => '0644',
    # The symlink target, when ensure => link
    target      => '/etc/httpd/httpd.conf'
    ,
    # Whether to recursively manage a directory (when ensure => directory)
    recurse     => true,
}
```

Executing commands

We can run plain commands using Puppet's `exec` type. Since Puppet applies it at every run, either the command can be safely run multiple times or we have to use one of the `creates`, `unless`, `onlyif`, `refreshonly` arguments to manage when to execute it.

```
exec { 'get_my_file':
  # (namevar) The command to execute
  command  => "wget http://mysite/myf
ile.tar.gz -O /tmp/myfile.tar.gz",
  # The search path for the command. M
ust exist when command is not absolute
  # Often set in Exec resource default
  s
  path      => '/sbin:/bin:/usr/sbin:/
/usr/bin',
  # A file created by the command. It
  if exists, the command is not executed
  creates   => '/tmp/myfile.tar.gz',
  # A command or an array of commands,
  if any of them returns an error
  # the command is not executed
  onlyif    => 'ls /tmp/myfile.tar.gz
&& false',
  # A command or an array of commands,
  if any of them returns an error
  # the command IS executed
  unless    => 'ls /tmp/myfile.tar.gz'
,
}
```

Managing users

Puppet has native types to manage users and groups, allowing easy addition, modification and removal. Here are the main arguments of the `user` type:

```
user { 'joe':
  # (namevar) The user name
  name      => 'joe',
  # The user's status: 'present', 'absent', 'role'
  ensure    => 'present',
  # The user's id
  uid       => '1001',
  # The user's primary group id
  gid       => '1001',
  # Eventual user's secondary groups (use array for many)
  groups    => [ 'admins' , 'developers' ],
  # The user's password. As it appears in /etc/shadow
  # Use single quotes to avoid unwanted evaluation of $* as variables
  password   => '$6$ZFS5JFFRZc$FFDSvPZS
SFGVdXD1He$',
  # Typical users' attributes
  shell      => '/bin/bash',
  home       => '/home/joe',
  mode       => '0644',
}
```

Puppet language style

[Puppet language](#) as an official Style Guide which defines recommended rules on how to write the code.

The standard de facto tool to check the code style is `puppet-lint`.

Badly formatted example:

```
file {
  '/etc/issue':
    content => "Welcome to $fqdn",
    ensure  => present,
    mode    => 644,
    group   => "root",
    path    => "${issue_file_path}",
}
```

Corrected style:

```
file { '/etc/issue':
  ensure  => present,
  content => "Welcome to ${fqdn}",
  mode    => 0644,
```

```
group => 'root',
path  => $issue_file_path,
}
```

Resources reference - Practice

Create a manifest that contains all the explored resources: package, service, file, exec, user.

Use Puppet Lint to verify its style.

Language reference - Overview

Referencing a resource

Managing resources ordering

Metaparameters

Conditionals and

Comparison operators

Nodes and classes inheritance

Resource references

In Puppet any resource is uniquely identified by its type and its name.
We can't have 2 resources of the same type with the same name in a catalog.

We have seen that we declare resources with a syntax like:

```
type { 'name':  
    arguments => values,  
}
```

When we need to reference to them in our code the syntax is like:

```
Type['name']
```

Some examples:

```
file { 'motd': ... }
apache::virtualhost { 'example42.com': .
...
exec { 'download_myapp': .... }
```

are referenced, respectively, with

```
File['motd']
Apache::Virtualhost['example42.com']
Exec['download_myapp']
```

Resource defaults

It's possible to set default argument values for a resource in order to reduce code duplication. The syntax is:

```
Type {  
    argument => value,  
}
```

Common examples:

```
Exec {  
    path => '/sbin:/bin:/usr/sbin:/usr/bin  
,  
}  
  
File {  
    mode  => 0644,  
    owner => 'root',  
    group => 'root',  
}
```

Resource defaults can be overridden when declaring a specific resource of the same type.

Note that the "Area of Effect" of resource defaults might bring unexpected results. The general suggestion is:

Place **global** resource defaults in /etc/puppet/manifests/site.pp outside any node definition.

Place **local** resource defaults at the beginning of a class that uses them (mostly for clarity sake, as they are parse-order independent).

Nodes inheritance

On the PuppetMaster we can define with the **node** definition the resources to apply to any node.

It is possible to have an inheritance structure for nodes, so that resources defined for a node are automatically included in an inheriting node.

An example:

```
node 'general' { ... }

node 'www01' inherits general { ... }
```

In Puppet versions prior to 3, it was possible to use nodes inheritance also to set variables and override them at different inheritance levels, and refer to these variables with their "short" name (not fully qualified). When using this approach it was important to avoid the

When using this approach it was important to avoid the inclusion on classes in the inheritance tree, since some variables could be evaluated in an unexpected way.

This is no more possible because variables are not more dynamically scoped, and generally speaking nodes inheritance has been deprecated.

Class inheritance

In Puppet classes are just containers of resources and have nothing to do with OOP classes. Therefore the meaning of class inheritance is somehow limited to few specific cases.

When using class inheritance, the main class ('puppet' in the sample below) is always evaluated first and all the variables and resource defaults it sets are available in the scope of the child class ('puppet::server').

Moreover the child class can override the arguments of a resource defined in the main class. Note the syntax used when referring to the existing resource File[/etc/puppet/puppet.conf]:

```
class puppet {
```

```
file { '/etc/puppet/puppet.conf':
  content => template('puppet/client/p
uppet.conf'),
}
}

class puppet::server inherits puppet {
  File['/etc/puppet/puppet.conf'] {
    content => template('puppet/server/p
uppet.conf'),
  }
}
```

Run Stages

In Puppet 2.6 it has been introduced the concept of Run Stages to help users in managing the order of dependencies when applying resources.

Puppet (> 2.6) provides a default **main** stage, we can

add any number or further stages with the stage resource type:

```
stage { 'pre':  
    before => Stage['main'],  
}
```

Which is equivalent to:

```
stage { 'pre': }  
Stage['pre'] -> Stage['main']
```

We can assign any class to a defined stage with the stage metaparameter:

```
class { 'yum':  
    stage => 'pre',  
}
```

Do not abuse of run stages (be vary of dependency cycles)!

[Official documentation on Run Stages](#)

Metaparameters

Metaparameters are parameters available to any resource type, they can be used for different purposes:

Manage dependencies (**before, require, subscribe, notify, stage**)

Manage resources' application policies (**audit, noop, schedule, loglevel**)

Add information to a resource (**alias, tag**)

Official documentation on Metaparameters

Managing dependencies

Puppet language is declarative and not procedural: it defines states, the order in which resources are written in manifests does not affect the order in which they are applied to the desired state.

To manage resources ordering, there are 3 different

methods, which can coexist:

1 - Use the metaparameters **before**, **require**, **notify**, **subscribe**

2 - Use the **Chaining arrows** (compared to the above metaparameters: **->** , **<-** , **<~** , **~>**)

3 - Use run stages

Managing dependencies - before | notify

In a typical Package/Service/Configuration file example we want the package to be installed first, configure it and then start the service, eventually managing its restart if the config file changes.

This can be expressed with metaparameters:

```
package { 'exim':
  before => File['exim.conf'],
}

file { 'exim.conf':
  notify => Service['exim'],
}

service { 'exim':
```

which is equivalent to

```
Package['exim'] -
> File['exim.conf'] ~> Service['exim']
```

Managing dependencies - require | subscribe

The previous example can be expressed using the alternative reverse metaparameters:

```
package { 'exim':  
}  
  
file { 'exim.conf':  
  require => Package['exim'],  
}  
  
service { 'exim':  
  subscribe => File['exim.conf'],  
}
```

which can also be expressed like:

```
Service['exim'] <~ File['exim.conf'] <-
  Package['exim']
```

Conditionals

Puppet provides different constructs to manage conditionals inside manifests:

Selectors allows to set the value of a variable or an argument inside a resource declaration according to the value of another variable.

Selectors therefore just returns values and are not used to manage conditionally entire blocks of code.

case statements are used to execute different blocks of code according to the values of a variable. It's possible to have a default block for unmatched entries.

Case statements are NOT used inside resource declarations.

if elif else conditionals, like case, are used to execute different blocks of code and can't be used inside resources declarations.

We can use any of Puppet's comparison expressions and we can combine more than one for complex patterns matching.

unless is somehow the opposite of **if**. It evaluates a boolean condition and if it's *false* it executes a block of code. It doesn't have *elsif / else* clauses.

Sample: Assign a variable value

Selector for variable's assignment

```
$package_name = $osfamily ? {
  'RedHat' => 'httpd',
  'Debian' => 'apache2',
  default   => undef,
}
```

case

```
case $::osfamily {
  'Debian': { $package_name = 'apache2' }
  'RedHat': { $package_name = 'httpd' }
  default: { notify { "Operating system $::operatingsystem not supported" } }
}
```

if elif else

```
if ::osfamily == 'Debian' {  
  $package_name = 'apache2'  
} elsif ::osfamily == 'RedHat' {  
  $package_name = 'httpd'  
} else {  
  notify { "Operating system ::operati  
ngsystem not supported" }  
}
```

Comparing strings:

operators

Puppet supports some common comparison operators:
== != < > <= >= =~ !~ and the somehow less common in

```
if $::osfamily == 'Debian' { [ ... ] }

if $::kernel != 'Linux' { [ ... ] }

if $::uptime_days > 365 { [ ... ] }

if $::operatingsystemrelease <= 6 { [ .. ] }
```

Comparing strings: in operator and combinations

in operator

The in operator checks if a string present in another string, an array or in the keys of an hash

```
if '64' in $::architecture
  if $monitor_tool in [ 'nagios' , 'icinga' , 'sensu' ]
```

expressions combinations

It's possible to combine multiple comparisons with **and** and **or**

```
if ($::osfamily == 'RedHat')
and ($::operatingsystemrelease == '5') {
    [ ... ]
}

if ($::osfamily == 'Debian') or ($::osfa-
mily == 'RedHat') {
    [ ... ]
}
```

Exported resources

When we need to provide to an host informations about resources present in another host, we need **exported resources**: resources declared in the catalog of a node (based on its facts and variables) but applied (collected) on another node.

Resources are declared with the special @@ notation which marks them as exported so that they are not applied to the node where they are declared:

```
@@host { $::fqdn:
  ip  => $::ipaddress,
}

@@concat::fragment { "balance-fe-
${::hostname}",
  target  => '/etc/haproxy/haproxy.cfg',
  content => "server ${::hostname} ${::i
paddress} maxconn 5000"
  tag      => "balance-fe",
}
```

Once a catalog containing exported resources has been applied on a node and stored by the PuppetMaster (typically on PuppetDB), the exported resources can be collected with the spaceshift syntax (where is possible to specify search queries):

```
Host << || >>
Concat::Fragment <<| tag == "balance-
fe" |>>
Sshkey <<| |>>
Nagios_service <<|||>>
```

Exported resources - Configuration

In order to use exported resources we need to enable on the Puppet Master the **storeconfigs** option and specify the backend to use.

We can do this configuring a PuppetMaster to use PuppetDB:

```
storeconfigs = true
storeconfigs_backend = puppetdb
```

In earlier Puppet versions storeconfigs is based on ActiveRecord, which is considerably slower.

Modules - Overview

Modules structure and conventions

Erb templates

Patterns for modules reusability

Principles of modules testing

Modules documentation

Modules

Self Contained and Distributable *recipes* contained in a directory with a predefined structure

Used to manage an application, system's resources, a local site or more complex structures

Modules must be placed in the Puppet Master's modulepath

```
puppet config print modulepath  
/etc/puppet/modules:/usr/share/puppet/modules
```

Puppet module tool to interface with [Puppet Modules](#)
Forge

```
puppet help module
[...]
ACTIONS:
  build      Build a module release package.
  changes    Show modified files of an installed module.
  generate   Generate boilerplate for a new module.
  install    Install a module from the Puppet Forge or an archive.
  list       List installed modules
  search     Search the Puppet Forge for a module.
  uninstall  Uninstall a puppet module
  .
  upgrade   Upgrade a puppet module.
```

GitHub, also, is full of [Puppet modules](#)

Paths of a module

Modules have a standard structure:

```
mysql/          # Main module directory  
mysql/manifests/ # Manifests directory.  
                  Puppet code here. Required.  
mysql/lib/       # Plugins directory. Ruby code here  
mysql/templates/ # ERB Templates directory  
mysql/files/     # Static files directory  
mysql/spec/      # Puppet-  
                  rspec test directory  
mysql/tests/     # Tests / Usage examples directory  
  
mysql/Modulefile # Module's metadata descriptor
```

This layout enables useful conventions

Modules paths conventions

Classes and defines autoloading:

```
include mysql
# Main mysql class is placed in: $modulepath/mysql/manifests/init.pp

include mysql::server
# This class is defined in: $modulepath/mysql/manifests/server.pp

mysql::conf { ...}
# This define is defined in: $modulepath/mysql/manifests/conf.pp

include mysql::server::ha
# This class is defined in: $modulepath/mysql/manifests/server/ha.pp
```

Provide files based on Erb Templates (Dynamic content)

```
content => template('mysql/my.cnf.erb'),
# Template is in: $modulepath/mysql/templates/my.cnf.erb
```

Provide static files (Static content). Note we can't use content AND source for the same file.

```
source => 'puppet:///modules/mysql/my.cnf'
# File is in: $modulepath/mysql/files/my.cnf
```

Erb templates

Files provisioned by Puppet can be Ruby ERB templates

In a template all the Puppet variables (facts or user assigned) can be used :

```
# File managed by Puppet on <%= @fqdn %>
search <%= @domain %>
```

But also more elaborated Ruby code

```
<% @dns_servers.each do |ns| %>
  nameserver <%= ns %>
<% end %>
```

The computed template content is placed directly inside the catalog
(Sourced files, instead, are retrieved from the puppetmaster during catalog application)

Principles behind a Reusable Module

Data Separation

- Configuration data is defined outside the module
- Module behavior can be managed via parameters
- Allow module's extension and override via external data

Reusability

- Support different OS. Easily allow new additions.
- Customize behavior without changing module code
- Do not force author idea on how configurations should be provided

Standardization

- Follow PuppetLabs style guidelines (puppet-lint)
- Have coherent, predictable and intuitive interfaces
- Provide contextual documentation (puppetdoc)

Interoperability

- Limit cross-module dependencies
- Allow easy modules cherry picking
- Be self contained, **do not interfere with other modules' resources**

Modules reusability patterns: Template + Options hash

Check this blog post for details.

Classes and defines should expose parameters that allow to override the used templates and set a custom hash of configurations.

```
class redis (
  $config_file_template = 'redis/redis.conf.erb',
  $options_hash         = {},
) {
  file { '/etc/redis/redis.conf':
    content => template($config_file_template),
  }
}
```

Template + Options hash example:

Given the previous class definition, we can configure it with this sample Hiera data, in YAML format:

```
---
redis::config_file_template: 'site/redis
/redis.conf.erb'
redis::options_hash:
  port: '12312'
  bind: '0.0.0.0'
  masterip: '10.0.42.50'
  masterport: '12350'
  slave: true
```

The referenced template stays in our site module, in `$modulepath/site/templates/redis/redis.conf.erb` and may look like:

```
port <%= @options_hash['port'] %>
bind <%= @options_hash['bind'] %>
<% if @options_hash['slave'] == true -%>
  slaveof <%= @options_hash['masterip'] %>
    <%= @options_hash['masterport'] %>
<% end -%>
```

Modules reusability patterns: Users override on included classes

Sometimes modules need some prerequisites or have to manage resources related to other applications or may provide the same functionality in different ways. Whenever this is needed, let users provide custom versions for the relevant classes.

This module manages the Puppet Master in a dedicated class and exposes a parameter that allows users to provide a custom class (or an alternative class eventually provided by the same module) instead of the default one.

```
class puppet {
  $server_class = '::puppet::server',
} {

  if $server_class {
    include $server_class
  }
}
```

Using Hiera we can override with:

```
---
puppet::server_class: '::site::puppet::haserver'
```

This implies that we need to create in \$modulepath/site/manifests/puppet/haserver.pp a class like:

```
class ::site::puppet::haserver {
  # The resources for our HA Puppet Master setup
}
```

Testing Modules

• unit code testing can be done at different levels

Puppet code testing can be done at different levels
with different tools

puppet parser validate <manifest.pp> - Checks the syntax of a manifest

puppet-lint <manifest.pp> - A gem that checks the style of a manifest

puppet-rspec - A gem that runs rspec unit tests on a module (Based on compiled catalog)

cucumber-puppet - A gem that runs cucumber tests a module (Based on compiled catalog) OBSOLETE

puppet-rspec-system - A gem that creates Vagrant VM and check for the expected results of real Puppet runs

Beaker - A gem that runs acceptance tests on multiple Vagrant VM

**Modules
documentation with
Puppet Doc**

```
$ puppet doc [--all] --mode rdoc [--outputdir ] [--debug|--verbose] [--trace]
  [--modulepath ] [--manifestdir ] [--config ]
```

Comment classes as below:

```
# Class: apache
#
# This module manages Apache
#
# Parameters:
#
# Actions:
#
# Requires:
#
# [Remember: No empty lines between comments and class definition]
class apache {
    ...
}
```

Hiera - Overview

Principles behind Hiera

Configuring hierarchies

Hiera plugins

**Using Hiera from Puppet and the
CLI**

Introduction to Hiera

Hiera is the **key/value lookup tool** of reference where to store Puppet user data.

It provides an highly customizable way to lookup for parameters values based on a custom hierarchy using many different backends for data storage.

It provides a command line tool `hiera` that we can use to interrogate directly the Hiera data and functions to be used inside Puppet manifests: `hiera()` , `hiera_array()` , `hiera_hash()` , `hiera_include()`

Hiera is installed by default with Puppet version 3 and

is available as separated download on earlier version
(Installation instructions).

We need Hiera only on the PuppetMaster (or on any node, if we have a masterless setup).

Currently version 1 is available, here is its official documentation.

Hiera configuration: **[hiera.yaml](#)**

Hiera's configuration file is in yaml format and is called `hiera.yaml` here we define the hierarchy we want to use and the backends where data is placed, with backend specific settings.

Hiera's configuration file path is different according to how it's invoked:

From the Command Line and Ruby code

Default path: `/etc/hiera.yaml`

From Puppet

Default path for Puppet OpenSource:
`/etc/puppet/hiera.yaml`

Default path for Puppet Enterprise:
`/etc/puppetlabs/puppet/hiera.yaml`

It's good practice the symlink these alternative configuration files in order to avoid inconsistencies when using Hiera from the shell line or within [Puppet](#) manifests:

```
ln -  
s /etc/hiera.yaml /etc/puppet/hiera.yaml
```

[Default configuration](#)

By default Hiera does not provide a configuration file.
The default settings are equivalent to this:

```
---  
:backends: yaml  
:yaml:  
  :datadir: /var/lib/hiera  
  :hierarchy: common  
  :logger: console
```


Hiera Backends

One powerful feature of Hiera is that the actual key-value data can be retrieved from different backends.

With the :backends global configuration we define which backends to use, then, for each used backend we can specify backend specific settings.

Build it backends:

yaml - Data is stored in yaml files (in the :datadir directory)

json - Data is stored in json files (in the :datadir directory)

puppet - Data is defined in Puppet (in the :datasouce class)

Extra backends:

Many additional backends are available, the most interesting ones are:

gpg - Data is stored in GPG encrypted yaml files

http - Data is retrieved from a REST service

mysql - Data is retrieved from a Mysql database

redis - Data is retrieved from a Redis database

Custom backends

It's relatively easy to write custom backends for Hiera. Here are some development instructions

Hierarchies

With the `:hierarchy` global setting we can define a string or an array of data sources which are checked in order, from top to bottom.

When the same key is present on different data sources by default is chosen the top one. We can override this setting with the `:merge_behavior` global configuration. Check this page for details.

In hierarchies we can interpolate variables with the `%{}` notation (variables interpolation is possible also in other parts of `hiera.yaml` and in the same data sources).

This is an example Hierarchy:

```
---  
:hierarchy:  
- "nodes/%{::clientcert}"  
- "roles/%{::role}"  
- "%{::osfamily}"  
- "%{::environment}"  
- common
```

Note that the referenced variables should be expressed with their fully qualified name. They are generally facts or Puppet's top scope variables (in the above example, `$::role` is not a standard fact, but is generally useful to have it (or a similar variable that identifies the kind of server) in our hierarchy).

If we have more backends, for each backend is evaluated the full hierarchy.

We can find some real world hierarchies samples in this [Puppet Users Group post](#)

More information on hierarchies [here](#).

Using Hiera in Puppet

The data stored in Hiera can be retrieved by the PuppetMaster while compiling the catalog using the `hiera()` function.

In our manifests we can have something like this:

```
$my_dns_servers = hiera("dns_servers")
```

Which assigns to the variable `$my_dns_servers` (can have any name) the top value retrieved by Hiera for the key `dns_servers`

We may prefer, in some cases, to retrieve all the values retrieved in the hierarchy's data sources of a given key and not the first use, use `hiera_array()` for that.

```
$my_dns_servers = hiera_array("dns_servers")
```

If we expect an hash as value for a given key we can use the `hiera()` function to retrieve the top value found or `hiera_hash` to merge all the found values in a single hash:

```
$openssh_settings = hiera_hash("openssh_settings")
```

Extra parameters

All these `hiera` functions may receive additional parameters:

Second argument: `default` value if no one is found

Third argument: `override` with a custom data source added at the top of the configured hierarchy

```
$my_dns_servers = hiera("dns_servers", "8.8.8", "$country")
```

Puppet 3 data bindings

With Puppet 3 Hiera is shipped directly with Puppet and an automatic hiera lookup is done for each class' parameter using the key **\$class::\$argument**: this functionality is called data bindings or automatic parameter lookup.

For example in a class definition like:

```
class openssh (
    template = undef,
) { . . . }
```

Puppet 3 automatically looks for the Hiera key `openssh::template` if no value is explicitly set when declaring the class.

To emulate a similar behaviour on pre Puppet 3 we should write something like:

```
class openssh (
    template = hiera("openssh::template"
),
) { . . . }
```

If a default value is set for an argument that value is used only when user has not explicitly declared value for that argument and Hiera automatic lookup for that argument doesn't return any value.

Using hiera from the command line

Hiera can be invoked via the command line to interrogate the given key's value:

```
hiera dns_servers
```

This will return the default value as no node's specific information is provided. More useful is to provide the whole facts' yaml of a node, so that the returned value can be based on the dynamic values of the hierarchy. On the Puppet Masters the facts of all the managed clients are collected in `Svardir.yaml!facts` so this is the best place to see how Hiera evaluates keys for different clients:

```
hiera dns_servers --  
yaml /var/lib/puppet/yaml/facts/<node>.y  
aml
```

We can also pass variables useful to test the hierarchy, directly from the command line:

```
hiera ntp_servers operatingsystem=Centos  
hiera ntp_servers operatingsystem=Ce  
ntos hostname=jenkins
```

To have a deeper insight of Hiera operations use the debug (-d) option:

```
hiera dns_servers -d
```

To make an hiera array lookup (equivalent to `hiera_array()`):

```
hiera dns_servers -a
```

To make an hiera hash lookup (equivalent to `hiera_hash()`):

```
hiera openssh::settings -h
```

Puppet architectures - Overview

The Components of a Puppet

architecture

Where to define classes

Where to define parameters

Where to place files

Puppet security

Components of a Puppet architecture

Tasks we deal with

Definition of the **classes** to be included in each node

Definition of the **parameters** to use for each node

Definition of the configuration **files** provided to the nodes

Components

/etc/puppet/manifests/site.pp - The default manifests loaded by the Master

ENC - The (optional) External Node Classifier

Idap - (Optional) LDAP backend

Hiera - Data key-value backend

Public modules - Public shared modules

Site modules - Local custom modules

Where to define classes

The classes to include in each node can be defined on:

/etc/puppet/manifests/site.pp - Top or Node scope variables

ENC - Under the classes key in the provided YAML

Idap - puppetClass attribute

Hiera - Via the `hiera_include()` function

Site modules - In roles and profiles or other grouping classes

Where to define parameters

The classes to include in each node can be defined on:

/etc/puppet/manifests/site.pp - Under the node statement

ENC - Following the ENC logic

Idap - puppetVar attribute

Hiera - Via the `hiera()`, `hiera_hash()`, `hiera_array()` functions of Puppet 3 Data Bindings

Shared modules - OS related settings

Site modules - Custom and logical settings

Facts - Facts calculated on the client

Where to define files

Shared modules - Default templates populated via

module's params

Site modules - All custom static files and templates

Hiera - Via the **hiera-file** plugin

Fileserver custom mount points

Code workflow management

Puppet code should stay under a SCM (Git, Subversion, Mercurial... whatever)

We must be able to test the code before committing to production

Testing Puppet code syntax is easy, testing its effects is not

Different testing methods for different purposes

Different code workflow options

Recommended: Couple Puppet environments with code branches More info

Puppet Security considerations

Client / Server communications are encrypted with SSL x509 certificates

By default the CA on the PuppetMaster requires manual signing of certificate requests

```
# In puppet.conf under [master] section:  
autosign = false
```

On the Puppet Master (as unprivileged user) runs a service that must be reached by every node (port 8140 TCP)

On the Client Puppet runs a root but does not expose a public service

Sensitive data might be present in the catalog and in reports

Sensitive data might be present in the puppet code (Propagated under a SCM)

Security advisories

PuppetDB - Overview

What PuppetDB does

Installation and configuration

API and query language

The puppetdbquery module

Introduction to PuppetDB

PuppetDB is a "fast, scalable and reliable" data warehouse for Puppet.

It saves data generated by Puppet: nodes' **facts**, **catalogs** and **reports**.

It can also work as a performant backend for **exported resources** being a recommended alternative to the older ActiveRecord storeconfigs interface.

It provides a faster backend also for Puppet's

Inventory Service.

PuppetDB is opensource, written in Clojure (it requires >= JDK 1.6 and PuppetMasters with Puppet >= 2.7.12), and is delivered as an autonomous software.

The best source from where to retrieve it are the official Puppetlabs repositories.

PuppetDB installation

For a complete overview of the installation process check the official documentation and choose between installatio from puppetdb module or from packages.

PuppetDB can persist data either on an embedded HSQLDB database or on PostgreSQL. The latter is definitively recommended on production environment where there are a few dozens of servers or more.

The configuration files are:

```/etc/puppetdb/conf.d/``` is the configuration directory, here we may have different .ini files where to configure \*\*[global]\*\* settings, \*\*[database]\*\* backends, \*\*[command-processing]\*\* options, \*\*[jetty]\*\* parameters for HTTP connections and \*\*[repl]\*\* settings for remote runtime configurations (used for development/debugging).

/etc/puppet/puppetdb.conf is the configuration file for **Puppet** with the settings to be used by the PuppetDB terminus

# **PuppetDB console and tools**

---

## **Integrated console: Anaylize PuppetDB performance**

PuppetDB provides a performances dashboard out of the box, we can use it to check how the software is working: <http://:8080/dashboard/>.

This is integral part of the PuppetDB software.

## **Puppet Board: Query PuppetDB from the web**

A nice frontend that allows interrogation of the PuppetDB from the web is PuppetBoard.

This software is contributed from the community.

## Puppetdbquery module: Query PuppetDB from Puppet manifests

An incredibly useful module that provides faces, functions and hiera backends that work with PuppetDB is the puppetdbquery module. We can install it also from the Forge:

```
puppet module install dalen-puppetdbquery
```

This software is contributed from the community (and is becoming a standard de facto for PuppetDB querying inside Puppet modules).

PuppetDB exposes an [HTTP API](#) that uses a Command/Query Responsibility Separation (CQRS) pattern:

#### **REST for reading**

A standard REST API is used to query data.

The current (API v3) available endpoints are: metrics, fact-names, facts, nodes, resources, reports, events, event-counts, aggregate-event-counts, server-time

#### **COMMANDS for writing**

Explicit **commands** are used (via HTTP using the /commands/ URL) used to populate and modify data.

The current **commands** are: replace catalog, replace facts, deactivate node, store report.

#### **API versions**

There are different versions of the APIs as they evolve with PuppetDB versions.

As October 2013, Version 1 of the API is deprecated, Version 2 and 3 (the latter adds new endpoints and is recommended) are both supported.

We can access a specific version of the [API](#) using the relevant prefix:

```
http[s]://puppetdb.server/v#/<endpoint>/
[NAME]/[VALUE][?query=<QUERY STRING>]
```

#### **Query language**

Queries to the REST endpoints can define search scope and limitations for the given endpoint. The query are sent as an "URL-encoded JSON array in prefix notation". Check the online tutorial for details.

# Using Puppetdbquery module

---

The PuppetDbQueryModule is developed by a community member, Erik Dalen, and is the most used and useful module available to work with PuppetDB: it provides command lines tools (as [Puppet Faces](#)), functions to query PuppetDB and a PuppetDB based Hiera backend.

## Query format

---

All the queries we can do with this module are in the format

```
Type[Name]
{attribute1=foo and attribute2=bar}
```

by default they are made on normal resources, use the

**@@** prefix to query exported resources.

The comparison operators are: **=**, **!=**, **>**, **<** and **~**

The expressions can be combined with **and**, **not** and **or**

## Using Puppetdbquery module

### Command line

The module introduces, as a Puppet face, the **query** command:

```
puppet help query
puppet query facts '(osfamily=RedHat and
operatingsystemversion=6)'
```

### Functions

The functions provided by the module can be used inside manifests to populate the catalog with data retrieved on PuppetDB.

**query\_nodes** takes 2 arguments: the query to use and (optional) the fact to return (by default it provides the certname). It returns an array:

```
$webservers = query_nodes('osfamily=Debi
an and Class[Apache]')
$webserver_ip = query_nodes('osfamily=De
bian and Class[Apache]', ipaddress)
```

**query\_facts** requires 2 arguments: the query to use to discover nodes and the list of facts to return for them. It returns a nested hash in JSON format.

```
query_facts('Class[Apache]
{port=443}', ['osfamily', 'ipaddress'])
```

# MCollective - Overview

---

## Essentials

### Installation and configuration

### The mco command

## MCollective essentials

---

An orchestration framework that allows massively parallel actions on the infrastructure's server

Based on 3 components:

One or more central consoles, from where the **mco** command can be issued

The infrastructure nodes, where an **agent** receives and executes the requested actions

A middleware **message broker**, that allows communication between master(s) and agents

Possible actions are based on plugins

Common plugins are: **service**, **package**, **puppetd**, **filemgr**...

Simple RCPs allow easy definitions of new plugins or actions.

Security is handled at different layers: transport, authentication and authorization via different plugins.  
Here is a Security Overview

# Installation

---

Some distros provide native mcollective packages but it's definitively recommended to use PuppetLabs repositories to have recent mcollective versions and packages for the needed dependencies.

On the nodes to be managed just install the **mcollective** package, it's configuration file is **/etc/mcollective/server.cfg**

On the administrative nodes, from where to control the infrastructure, install the **mcollective-client** package, it's configuration file is **/etc/mcollective/client.cfg**

On both clients and servers we have to install the agent plugins

On the middleware server(s) install **activemq** or **rabbitmq**

# Configuration

---

Configuration files are similar for server and client.

They are in a typical setting = value format and # are used for comments.

We have to define connection settings to the middleware, the settings for the chosen security plugins, the facts source and other general settings.

The setup of the whole infrastructure involves some parameters to be consistent on the mcollective config files and the message broker ones (credentials and eventually certificates).

There can be different setups, according to the security plugins used, the recommended one is well described in the Standard Mcollective Deployment

# Using the mco command

---

The **mco** command is installed on the Mcollective clients and can be used to perform actions on any node of the infrastructure.

General usage:

```
mco [subcommand] [options] [filters]
mco rpc [options] [filters] <agent> <action> [<key=val> <key=val>]
```

Most used subcommands are **ping** , **rpc** , **inventory**, **find**

Most used options are: **-batch SIZE** (send requests in batches), **-q** (quiet), **-j** (produce Json output), **-v** (verbose)

Most used filters are: **-F <fact=val>** (Match hosts with the specified fact value), **-I** (Match host with the given name).

Sample commands:

```
mco help
mco help rpc
mco ping
mco rpc service status service=mysql
mco ping -F datacenter=eu-1
```

# Puppet reporting - Overview

---

**Reporting overview**

**Understanding puppet run output**

## Reporting

---

In reports Puppet stores information of what has changed on the system after a Puppet run

Reports are sent from the client to the Master, if report is enabled

```
On client's puppet.conf [agent]
report = true
```

On the Master different report processors may be enabled

```
On server's puppet.conf [master]
reports = log,tagmail,store,https
reporturl = http://localhost:3000/reports
```

## Understanding Puppet runs output

Puppet is very informative about what it does on the functions

```
/Stage[main]/Example42::CommonSetup/File
[/etc/motd]/content:
--- /etc/motd 2012-12-
13 10:37:29.000000000 +0100
+++ /tmp/puppet+file20121213-19295-qe4wv2-0 2012-12-
13 10:38:19.000000000 +0100
@@ -4,4 +4,4 @@
-
Last Puppet Run: Thu Dec 13 10:36:58 CET
2012
+Last Puppet Run: Thu Dec 13 10:38:00 CET
2012
Info: /Stage[main]/Example42::CommonSetup/File[/etc/motd]: Filebucketed /etc/motd to main with sum 623bcd5f8785251e2cd00a88438d6d08
/Stage[main]/Example42::CommonSetup/File
[/etc/motd]/content: content changed '{md5}623bcd5f8785251e2cd00a88438d6d08' to '{md5}91855575e2252c38ce01efedf1f48cd3'
```

The above lines refer to a change made by Puppet on /etc/motd  
The File[/etc/motd] resource is defined in the class example42::commonsetup  
A diff is shown of what has changed (the diff appears when running puppet agent -t in interactive mode)

The original file has been filebucketed (saved) with  
checksum 623bcd5f8785251e2cd00a88438d6d08  
We can retrieve the original file in  
`/var/lib/puppet/clientbucket/6/2/3/b/c/d/5/f/623bcd5f8785251e2cd00a88438d6d08/contents`  
We can search for all the old versions of /etc/motd with

```
grep -R '/etc/motd' /var/lib/puppet/clientbucket/*
```

# Developing Puppet extensions - Overview

---

**Developing custom facts**

**Developing custom types and provides**

**Developing functions**

## Developing Facts

---

Facts are generated on the client before the evaluation of the Puppet code on the server.

Facts provide variables that can be used in Puppet code and templates.

A simple custom type that just executes a shell command:

```
require 'facter'
```

```
require 'facter'

Facter.add("last_run") do
 setcode do
 Facter::Util::Resolution.exec('date')
 end
end
```

This file should be placed in  
`/lib/facter/acpi_available.rb`

## Developing Types

Resource types can be defined in [Puppet](#) language or in Ruby as plugins.

Ruby types require one or more **providers** which manage low-level interaction with the underlying OS to provide the more abstract resource defined in Types.

An example of a type with many providers is package, which has more than 20 providers that manage packages on different OS.

A sample type structure (A type called `vcsrepo`, must have a path like `/lib/puppet/type/vcsrepo.rb`)

```
require 'pathname'

Puppet::Type.newtype(:vcsrepo) do
 desc "A local version control repository"

 feature :gzip_compression,
 "The provider supports explicit GZip compression levels"
 [...] List of the features

 ensurable do
 newvalue :present do
 provider.create
 end
 [...] List of the accepted values
 newparam(:source) do
 desc "The source URI for the repository"
 end
 [...] List of the accepted parameters
end
```



## Developing Providers

The vcsrepo type seen before has 5 different providers for different source control tools.

Here is, as an example, the git provider for the vcsrepo type (Code from Puppet Labs):

This is the full provider:

[lib/puppet/provider/vcsrepo/git.rb](#)

```
require File.join(File.dirname(__FILE__),
 '..', 'vcsrepo')

Puppet::Type.type(:vcsrepo).provide(:git)
 parent = Puppet::Provider::Vcsrepo
do
 desc "Supports Git repositories"
 commands git = 'git'
 defaultfor git => exists
 has_features :bare_repositories, :reference_tracking

 def create
 if @resource.value(:source)
 init_repository(@resource.value(:path))
 else
 clone_repository(@resource.value(:source), @resource.value(:path))
 if @resource.value(:revision)
 if @resource.value(:ensure) == :bare
 notice "Ignoring revision for bare repository"
 else
 checkout_branch_or_reset
 end
 if @resource.value(:ensure) != :bare
 update_submodules
 end
 end
 end
 end

 def destroy
 FileUtils.rm_rf(@resource.value(:path))
 end
 [...]
end
```



# Developing Functions

Functions are Ruby code that is executed during compilation on the [Puppet Master](#).

They are used to interface with external tools, provide debugging or interpolate strings.

Imports parts of the [Puppet](#) language like include and template are implemented as functions.

```
module Puppet::Parser::Functions
 newfunction(:get_magicvar, :type => :r
 value, :doc => <<-EOS
This returns the value of the input variable. For example if you input role
it returns the value of '$role'.
EOS
) do |arguments|
 raise(Puppet::ParseError, "get_magicvar(): Wrong number of arguments " +
 "given (#
{arguments.size} for 1)") if arguments.size < 1

 my_var = arguments[0]
 result = lookupvar("#{my_var}")
 result = 'all' if (result == :undefined || result == '')
 return result
end
end
```

This file is placed in  
puppi/lib/puppet/parser/functions/get\_magicvar.rb.

# Scaling Puppet

---

**Optimizing code for performance**

**Optimize code for  
performance**

---

Reduce the number of resources per node  
For each resource Puppet has to serialize, deserialize,  
apply, report...  
Avoid use of resources for too many entities (Do

manage hundreds of users with the User type)  
Limit overhead of containing classes  
(A module with the openssh::package,  
openssh::service, openssh::configuration subclasses  
pattern is uselessly filled with extra resources)

Do not use the file type to deliver too large files or  
binaries:

For each file Puppet has to make a checksum to verify  
if it has changed

With source => , the content of the file is retrieved with  
a new connection to the PuppetMaster

With content => template() , the content is placed  
inside the catalog.

Avoid too many elements in a source array for file  
retrieval:

```
source => ["site/openssh/sshd.conf--
$::hostname ,
 "site/openssh/sshd.conf--
$environment-$role ,
 "site/openssh/sshd.conf-
$role ,
 "site/openssh/sshd.conf],
```

This checks 3 files (and eventually gets 3 404 errors  
from server) before getting the default ones.

## **Reduce PuppetMaster(s) load**

Run Puppet less frequently (default 30 mins)

Evaluate centrally managed, on-demand, Puppet Runs  
(ie: via MCollective)

Evaluate Master-less setups

Use a setup with Passenger to have multiple  
PuppetMasters childs on different CPUs

Disable Store Configs if not used

```
storeconfigs = false
```

Alternatively enable Thin Store Configs and Mysql  
backend

```
storeconfigs = true
thin_storeconfigs = true
dbadapter = mysql
dbname = puppet
dbserver = mysql.example42.com
dbuser = puppet
dbpassword = <%= scope.lookupvar('secret
::puppet_db_password') %>
```

Evaluate PuppetDB as backend (faster)

```
storeconfigs = true
storeconfigs_backend = puppetdb
```

